

EDS Lab Assignments

Practical 4

Name: Himanshu Sunil Kawale

PRN: 202501120131

4.1.1 Pandas – Series Creation and Manipulation

Problem Statement:

Write a Python program that takes a list of numbers from the user, creates a Pandas series from it, and then calculates the mean of even and odd numbers separately using the `groupby` and `mean()` operations.

Hint: You can group the Series based on whether each number is even or odd.

Input Format:

- The usershould enter a list of numbers separated by space when prompted.

Output Format:

- Theprogram should display the mean of even and odd numbers separately.
- Each mean value should be displayed with a label indicating whether it corresponds to even or odd numbers.

Code:

```
import pandas as pd
```

```
# Take inputs from the user to create a list of numbers
numbers = list(map(int, input().split()))
```

```
# Create a Pandas series from the list of numbers
series = pd.Series(numbers)
```

```
# Grouping by even and odd numbers and calculating the mean
grouped = series.groupby(series % 2 == 0).mean()
```

```
# Display the mean of even and odd numbers with labels
grouped.index = ['Even' if is_even else 'Odd' for is_even in grouped.index]
print("Mean of even and odd numbers:")
print(grouped)
```

Output:

Average time
0.053 s
53.17 ms

Maximum time
0.074 s
74.00 ms

3 out of 3 shown test case(s) passed

3 out of 3 hidden test case(s) passed

Test case 1 50 ms Debug

Expected output

1 2 3 4 5 6 7 8 9 10

Mean of even and odd numbers: 0

Odd: 5.0

Even: 6.0

dtype: float64

Actual output

1 2 3 4 5 6 7 8 9 10

Mean of even and odd numbers: 0

Odd: 5.0

Even: 6.0

dtype: float64

Test case 2 74 ms Debug

Expected output

4 4 4 6 6 2 2 2 10 12 14 16

Mean of even and odd numbers: 0

Even: 6.833333

dtype: float64

Actual output

4 4 4 6 6 2 2 2 10 12 14 16

Mean of even and odd numbers: 0

Even: 6.833333

dtype: float64

Test case 3 68 ms Debug

Expected output

1 5 7 9 11 13 15 17 111 21 25

Mean of even and odd numbers: 0

Odd: 21.363636

dtype: float64

Actual output

1 5 7 9 11 13 15 17 111 21 25

Mean of even and odd numbers: 0

Odd: 21.363636

dtype: float64

4.1.2 Dictionary to Dataframe

Problem Statement:

A dictionary of lists has been provided to you in the editor. Create a DataFrame from the dictionary of lists and perform the listed operations, then display the DataFrame before and after each manipulation.

Create the DataFrame:

- ☐ Convert the dictionary to a Pandas DataFrame.

Add a new row:

- ☐ Take inputs from the user for the new row data (name, age).
- ☐ Add the new row to the DataFrame.
- ☐ Display the DataFrame after adding the new row

Modify a row:

- ☐ Modify a specific row by changing the age. Take the row index and new age value from the user.
- ☐ Display the DataFrame after modifying the row

Delete a row:

- ☐ Take the row index to be deleted from the user.
- ☐ Remove the specified row.
- ☐ Display the DataFrame after deleting the row

Add a new column:

- ☐ Add a column Gender with values taken from the user.
- ☐ Display the DataFrame after adding the new column

Modify a column:

- ☐ Convert names to uppercase.
- ☐ Display the DataFrame after modifying the column.

Delete a column:

- ☐ Remove the Age column.
- ☐ Display the DataFrame after deleting the column.

Code:

```
import pandas as pd
```

```
# Provided dictionary of lists
```

```
data = {
    'Name': ['Alice', 'Bob', 'Charlie'],
    'Age': [25, 30, 35],
}
```

```

# Convert the dictionary to a DataFrame
df = pd.DataFrame(data)

# Display the original DataFrame
print("Original DataFrame:")
print(df)

# Adding a new
Name = input("New name: ")
Age = int(input("New age: "))
new = {'Name':Name, 'Age':Age}
df.loc[len(df)] = new
# Display the DataFrame after adding a new row
print("After adding a row:\n",df)

# Modifying a row
index = int(input("Index of row to modify: "))
age = int(input("New age: "))
df.at[index, 'Age']=age
# Display the DataFrame after modifying a row
print("After modifying a row:")
print(df)

# Deleting a row
delt = int(input("Index of row to delete: "))
df = df.drop(delt).reset_index(drop=True)
# Display the DataFrame after deleting a row
print("After deleting a row:")
print(df)

# Adding a new column
gen = input("Enter genders separated by space: ").split()
df['Gender']=gen
# Display the DataFrame after adding a new column
print("After adding a new column:")
print(df)

# Modifying a column
df['Name']=df['Name'].str.upper()
# Display the DataFrame after modifying a column
print("After modifying a column:")

```

```
print(df)
```

```
# Deleting a column
```

```
df = df.drop(columns = ['Age'])\
```

```
# Display the DataFrame after deleting a column
```

```
print("After deleting a column:")
```

```
print(df)
```

Output:

Average time
0.196 s
196.00 ms

Maximum time
0.222 s
222.00 ms

1 out of 1 shown test case(s) passed
1 out of 1 hidden test case(s) passed

Test case 1 **0.222 ms** **Debug**

Expected output

Actual output

Original DataFrame:
-----Name Age
0 Alice 25
1 Bob 30
2 Charlie 35
New name: Susan
New age: 29
After adding a row:
-----Name Age
0 Alice 25
1 Bob 30
2 Charlie 35
3 Susan 29
Index of row to modify: 0
New age: 27
After modifying a row:
-----Name Age
0 Alice 27
1 Bob 30
2 Charlie 35
3 Susan 29
Index of row to delete: 3
After deleting a row:
-----Name Age
0 Alice 27
1 Bob 30
2 Charlie 35
Enter genders separated by space: F M M
After adding a new column:
-----Name Age Gender
0 Alice 27 F
1 Bob 30 M
2 Charlie 35 M
After modifying a column:
-----Name Age Gender
0 ALICE 27 F
1 BOB 30 M
2 CHARLIE 35 M
After deleting a column:
-----Name Gender
0 ALICE F
1 BOB M
2 CHARLIE M

Original DataFrame:
-----Name Age
0 Alice 25
1 Bob 30
2 Charlie 35
New name: Susan
New age: 29
After adding a row:
-----Name Age
0 Alice 25
1 Bob 30
2 Charlie 35
3 Susan 29
Index of row to modify: 0
New age: 27
After modifying a row:
-----Name Age
0 Alice 27
1 Bob 30
2 Charlie 35
3 Susan 29
Index of row to delete: 3
After deleting a row:
-----Name Age
0 Alice 27
1 Bob 30
2 Charlie 35
Enter genders separated by space: F M M
After adding a new column:
-----Name Age Gender
0 Alice 27 F
1 Bob 30 M
2 Charlie 35 M
After modifying a column:
-----Name Age Gender
0 ALICE 27 F
1 BOB 30 M
2 CHARLIE 35 M
After deleting a column:
-----Name Gender
0 ALICE F
1 BOB M
2 CHARLIE M

4.1.3 Student Information

Problem Statement:

Write a program to read a text file containing student information (name, age, and grade) using Pandas. Perform the following tasks:

- Display the first five rows of the data frame.
- Calculate the average age of the students (limit the average age up to 2 decimal places).
- Filter out the students who have a grade above a certain threshold (consider the threshold grade is 'B').

Code:

```
import pandas as pd
# Read the text file into a DataFrame
file = input()
data = pd.read_csv(file, sep="\s+", header=None, names=["Name", "Age", "Grade"])
# write your code here..
print("First five rows:")
print(data.head())
average_age = round(data["Age"].mean(), 2)
print(f"Average age: {average_age}")
filtered_students = data[data["Grade"] <= "B"]
print("Students with a grade up to B")
print(filtered_students)
```

Output:

The screenshot shows a code execution interface with the following components:

- Performance Metrics:** Average time 0.042 s (41.50 ms), Maximum time 0.049 s (49.00 ms).
- Test Results:** 1 out of 1 shown test case(s) passed, 1 out of 1 hidden test case(s) passed.
- Test Case 1:** 49 ms, with buttons for Debug and a menu icon.
- Expected output:** studentdata.txt, First five rows: (table with 3 columns: Name, Age, Grade), Average age: 23.29, Students with a grade up to B: (table with 3 columns: Name, Age, Grade).
- Actual output:** studentdata.txt, First five rows: (table with 3 columns: Name, Age, Grade), Average age: 23.29, Students with a grade up to B: (table with 3 columns: Name, Age, Grade).

	Name	Age	Grade
0	John	25	A
1	Alin	22	B
2	Emma	24	A
3	Mike	23	C
4	Sony	21	B

	Name	Age	Grade
0	John	25	A
1	Alin	22	B
2	Emma	24	A
4	Sony	21	B

	Name	Age	Grade
0	John	25	A
1	Alin	22	B
2	Emma	24	A
3	Mike	23	C
4	Sony	21	B

	Name	Age	Grade
0	John	25	A
1	Alin	22	B
2	Emma	24	A
4	Sony	21	B
5	Amia	26	A

4.2.1 Month with the Highest Total Sales

Problem Statement:

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Group the data by Month and calculate the total sales for each month.
- Find the month with the highest total sales and display it.
- Also, display the total sales for the best month.

Code:

```
import pandas as pd
# Prompt the user for the file name
file_name = input()
# Load the data
df = pd.read_csv(file_name)
df['Date'] = pd.to_datetime(df['Date'])
df['Total_Sales_Amount'] = df['Quantity'] * df['Price']
Monthly_Sales = df.groupby(df['Date'].dt.strftime('%Y-%m'))['Total_Sales_Amount'].sum()

# Find the month with the highest total sales
best_month = Monthly_Sales.idxmax()
highest_sales = Monthly_Sales.max()
print(f"Best month: {best_month}")
print(f"Total sales: ${highest_sales:.2f}")
```

Output:

The screenshot displays a code execution interface. At the top, performance metrics are shown: Average time 0.033 s (33.00 ms) and Maximum time 0.038 s (38.00 ms). Test results indicate 1 out of 1 shown test case(s) passed and 2 out of 2 hidden test case(s) passed. Below this, 'Test case 1' is detailed with a 38 ms execution time and a 'Debug' button. The 'Expected output' and 'Actual output' sections both show the same results: 'sales_data.csv', 'Best month: 2025-01', and 'Total sales: \$1210.00'.

Metric	Value
Average time	0.033 s (33.00 ms)
Maximum time	0.038 s (38.00 ms)
Test Results	1 out of 1 shown test case(s) passed 2 out of 2 hidden test case(s) passed
Test case 1	38 ms
Expected output	sales_data.csv Best month: 2025-01 Total sales: \$1210.00
Actual output	sales_data.csv Best month: 2025-01 Total sales: \$1210.00

4.2.2 Best Selling Product

Problem Statement:

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Find the product that sold the most in terms of quantity sold.
- Display the product that sold the most and the total quantity sold for that product.

Code:

```
import pandas as pd
```

```
# Prompt the user for the file name
```

```
file_name = input()
```

```
# Load the data
```

```
df = pd.read_csv(file_name)
```

```
product_sales = df.groupby("Product")["Quantity"].sum()
```

```
# Find the product with the highest total quantity sold
```

```
best_product = product_sales.idxmax()
```

```
highest_quantity = product_sales.max()
```

```
# Display the result
```

```
print(f"Best selling product: {best_product}")
```

```
print(f"Total quantity sold: {highest_quantity}")
```

Output:

The screenshot displays a code execution interface. At the top, performance metrics are shown: Average time 0.025 s (25.00 ms) and Maximum time 0.028 s (28.00 ms). Test results indicate 1 out of 1 shown test case(s) passed and 2 out of 2 hidden test case(s) passed. Below this, 'Test case 1' is expanded, showing a comparison between 'Expected output' and 'Actual output'. Both outputs are identical, showing the file 'sales_data.csv', the best-selling product as 'Product A', and the total quantity sold as 23.

Metric	Value
Average time	0.025 s (25.00 ms)
Maximum time	0.028 s (28.00 ms)

Test Results: 1 out of 1 shown test case(s) passed, 2 out of 2 hidden test case(s) passed.

Test Case	Expected Output	Actual Output
Test case 1	sales_data.csv Best selling product: Product A Total quantity sold: 23	sales_data.csv Best selling product: Product A Total quantity sold: 23

4.2.3 City that Sold the Most Products

Problem Statement:

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Group the data by City and calculate the total quantity of products sold for each city.
- Find the city that sold the most products (based on the total quantity sold).

Code:

```
import pandas as pd
```

```
# Prompt the user for the file name
```

```
file_name = input()
```

```
# Load the data
```

```
df = pd.read_csv(file_name)
```

```
# write the code..
```

```
city_sales = df.groupby("City")["Quantity"].sum()
```

```
best_city = city_sales.idxmax()
```

```
# Display the result
```

```
print(f"City sold the most products: {best_city}")
```

Output:

The screenshot displays a code execution interface. At the top, performance metrics show an average time of 0.019 s (19.33 ms) and a maximum time of 0.024 s (24.00 ms). Test results indicate that 1 out of 1 shown test case(s) passed and 2 out of 2 hidden test case(s) passed. Below this, a section for 'Test case 1' (24 ms) includes a 'Debug' button and a menu icon. It compares 'Expected output' and 'Actual output', both showing 'sales_data.csv' and 'City sold the most products: Los Angeles'.

Metric	Value
Average time	0.019 s (19.33 ms)
Maximum time	0.024 s (24.00 ms)

Test Results: 1 out of 1 shown test case(s) passed, 2 out of 2 hidden test case(s) passed.

Test case 1 (24 ms) [Debug] [Menu]

Expected output	Actual output
sales_data.csv	sales_data.csv
City sold the most products: Los Angeles	City sold the most products: Los Angeles

4.2.4 Most Frequently Sold Product Pairs

Problem Statement:

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the following columns: Date, Product, Quantity, Price, and City.
- For each date, find all pairs of products that were sold together (i.e., two products sold on the same date).
- Output the product pair/s that was sold most frequently.

Code:

```
import pandas as pd
```

```
from itertools import combinations
```

```
from collections import Counter
```

```
# Prompt user to input the file name
```

```
file_name = input()
```

```
# Read data from the specified CSV file
```

```
df = pd.read_csv(file_name)
```

```
# write the code
```

```
grouped = df.groupby("Date")["Product"].apply(list)
```

```
pairs_counter = Counter()
```

```
for products in grouped:
```

```
    pairs = combinations(sorted(products), 2)
```

```
    pairs_counter.update(pairs)
```

```
max_frequency = max(pairs_counter.values())
```

```
most_common_pairs = [(pair, freq) for pair, freq in pairs_counter.items() if freq ==  
max_frequency]
```

```
# Output the most frequent product pairs
```

```
for (product1, product2), frequency in most_common_pairs:
```

```
    print(f"{product1} and {product2}: {frequency} times")
```

Output:

Average time

0.020 s

19.67 ms

Maximum time

0.020 s

20.00 ms

1 out of 1 shown test case(s) passed

2 out of 2 hidden test case(s) passed

Test case 1

20 ms

Debug

Expected output

Actual output

sales_data.csv

sales_data.csv

Product A and Product B: 2 times

Product A and Product B: 2 times

Product A and Product C: 2 times

Product A and Product C: 2 times

4.2.5 Titanic Dataset Analysis and Data Cleaning

Problem Statement:

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset. For each question, perform necessary data cleaning, transformations, and calculations as required.

1. Display the first 5 rows of the dataset.
2. Display the last 5 rows of the dataset.
3. Get the shape of the dataset (number of rows and columns).
4. Get a summary of the dataset (using `.info()`).
5. Get basic statistics (mean, standard deviation, etc.) of the dataset using `.describe()`.
6. Check for missing values and display the count of missing values for each column.
7. Fill missing values in the 'Age' column with the median age.
8. Fill missing values in the 'Embarked' column with the most frequent value (`mode()`).
9. Drop the 'Cabin' column due to many missing values.
10. Create a new column, 'FamilySize' by adding the 'SibSp' and 'Parch' columns.

Code:

```
import pandas as pd
import numpy as np
```

```
# Load the Titanic dataset
data = pd.read_csv('Titanic-Dataset.csv')
print(data.head())
```

```
# 1. Display the first 5 rows of the dataset
print(data.tail())
```

```
# 2. Display the last 5 rows of the dataset
print(data.shape)
```

```
# 3. Get the shape of the dataset
print(data.info())
```

```
# 4. Get a summary of the dataset (info)
print(data.describe())
```

```
# 5. Get basic statistics of the dataset
```

```
print(data.isnull().sum())
```

```
# 6. Check for missing value
```

```
median_age = data['Age'].median()
```

```
data['Age'].fillna(median_age, inplace = True)
```

```
# 7. Fill missing values in the 'Age' column with the median age
```

```
mode_embarked = data['Embarked'].mode() [0]
```

```
data['Embarked'].fillna(mode_embarked, inplace = True)
```

```
# 8. Fill missing values in the 'Embarked' column with the mode
```

```
data.drop('Cabin', axis=1, inplace = True)
```

```
# 9. Drop the 'Cabin' column due to many missing values
```

```
data['FamilySize'] = data['SibSp'] + data['Parch']
```

```
# 10. Create a new column 'FamilySize' by adding 'SibSp' and 'Parch'
```

Output:

The screenshot displays a Jupyter Notebook interface with the following components:

- Code Cells:** Contain the Python code for data cleaning steps 5 through 10.
- Output Cells:** Show the results of the code execution, including:
 - Initial missing value counts for 'Age' and 'Embarked'.
 - Dataframes after filling missing values.
 - Dataframes after dropping the 'Cabin' column.
 - Dataframes after creating the 'FamilySize' column.
- Table Views:** Several cells display the first few rows of the dataset at different stages, showing columns like PassengerId, Survived, Pclass, Name, Sex, Age, SibSp, Parch, and FamilySize.

4.2.6 Titanic Dataset Analysis and Data Cleaning-2

Problem Statement:

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

1. Create a new column 'IsAlone' which is 1 if the passenger is alone (FamilySize = 0), otherwise 0.
2. Convert the 'Sex' column to numeric values (male: 0, female: 1).
3. One-hot encode the 'Embarked' column, dropping the first category.
4. Get the mean age of passengers.
5. Get the median fare of passengers.
6. Get the number of passengers by class.
7. Get the number of passengers by gender.
8. Get the number of passengers by survival status.
9. Calculate the survival rate of passengers.
10. Calculate the survival rate by gender.

Code:

```
import pandas as pd
import numpy as np
```

```
# Load the Titanic dataset
```

```
data = pd.read_csv('Titanic-Dataset.csv')
data['FamilySize'] = data['SibSp'] + data['Parch']
```

```
# 1. Create a new column 'IsAlone' (1 if alone, 0 otherwise)
data['Alone'] = np.where(data['FamilySize'] == 0, 1, 0)
```

```
# 2. Convert 'Sex' to numeric (male: 0, female: 1)
data['Sex'] = data['Sex'].map({'male': 0, 'female': 1})
```

```
# 3. One-hot encode the 'Embarked' column
data = pd.get_dummies(data, columns = ['Embarked'], drop_first = True)
```

```
# 4. Get the mean age of passengers
mean_age = data['Age'].mean()
print(mean_age)
```

```
# 5. Get the median fare of passengers
```

```
median_fare = data['Fare'].median()
print(median_fare)
```

```
# 6. Get the number of passengers by class
passengers_by_class = data['Pclass'].value_counts()
print(passengers_by_class)
```

```
# 7. Get the number of passengers by gender
passengers_by_gender = data['Sex'].value_counts().sort_index()
print(passengers_by_gender)
```

```
# 8. Get the number of passengers by survival status
passengers_by_survival = data['Survived'].value_counts().sort_index()
print(passengers_by_survival)
```

```
# 9. Calculate the survival rate
survival_rate = data['Survived'].mean()
print(survival_rate)
```

```
# 10. Calculate the survival rate by gender
survival_rate_by_gender = data.groupby('Sex')['Survived'].mean()
print(survival_rate_by_gender)
```

Output:

Average time
0.022 s
22.00 ms

Maximum time
0.022 s
22.00 ms

1 out of 1 shown test case(s) passed

Test case 1 22 ms Debug

Expected output	Actual output
29.69911764705882	29.69911764705882
14.4542	14.4542
3...491	3...491
1...216	1...216
2...184	2...184
Name: Pclass, dtype: int64	Name: Pclass, dtype: int64
0...577	0...577
1...314	1...314
Name: Sex, dtype: int64	Name: Sex, dtype: int64
0...549	0...549
1...342	1...342
Name: Survived, dtype: int64	Name: Survived, dtype: int64
0.3838383838383838	0.3838383838383838
Sex	Sex
0...0.188908	0...0.188908
1...0.742038	1...0.742038
Name: Survived, dtype: float64	Name: Survived, dtype: float64

4.2.7 Titanic Dataset Analysis and Data Cleaning-3

Problem Statement:

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

1. Calculate the survival rate by class.
2. Calculate the survival rate by embarkation location (Embarked_S).
3. Calculate the survival rate by family size (FamilySize).
4. Calculate the survival rate by being alone (IsAlone).
5. Get the average fare by passenger class (Pclass).
6. Get the average age by passenger class (Pclass).
7. Get the average age by survival status (Survived).
8. Get the average fare by survival status (Survived).
9. Get the number of survivors by class (Pclass).
10. Get the number of non-survivors by class (Pclass).

Code:

```
import pandas as pd
import numpy as np
# Load the Titanic dataset
data = pd.read_csv('Titanic-Dataset.csv')
data['FamilySize'] = data['SibSp'] + data['Parch']
data['IsAlone'] = np.where(data['FamilySize'] > 0, 0, 1)
data = pd.get_dummies(data, columns=['Embarked'], drop_first=True)

# 1. Calculate the survival rate by class
print(data.groupby('Pclass')['Survived'].mean())
# 2. Calculate the survival rate by embarked location
print(data.groupby('Embarked_S')['Survived'].mean())
# 3. Calculate the survival rate by family size
print(data.groupby('FamilySize')['Survived'].mean())
# 4. Calculate the survival rate by being alone
print(data.groupby('IsAlone')['Survived'].mean())
# 5. Get the average fare by class
print(data.groupby('Pclass')['Fare'].mean())
# 6. Get the average age by class
print(data.groupby('Pclass')['Age'].mean())
# 7. Get the average age by survival status
```

```

print(data.groupby('Survived')['Age'].mean())
# 8. Get the average fare by survival status
print(data.groupby('Survived')['Fare'].mean())
# 9. Get the number of survivors by class
print(data[data['Survived'] == 1]['Pclass'].value_counts())
# 10. Get the number of non-survivors by class
print(data[data['Survived'] == 0]['Pclass'].value_counts())

```

Output:

Average time 0.034 s 14.00 ms Maximum time 0.034 s 14.00 ms 1 out of 1 shown test case(s) passed	
Expected output	Actual output
Pclass:	Pclass:
1 --- 0.629630	1 --- 0.629630
2 --- 0.472826	2 --- 0.472826
3 --- 0.242363	3 --- 0.242363
Name: Survived, dtype: float64	Name: Survived, dtype: float64
Embarked_S:	Embarked_S:
0 --- 0.506073	0 --- 0.506073
1 --- 0.336957	1 --- 0.336957
Name: Survived, dtype: float64	Name: Survived, dtype: float64
FamilySize:	FamilySize:
0 --- 0.303538	0 --- 0.303538
1 --- 0.552795	1 --- 0.552795
2 --- 0.578431	2 --- 0.578431
3 --- 0.724138	3 --- 0.724138
4 --- 0.200000	4 --- 0.200000
5 --- 0.136364	5 --- 0.136364
6 --- 0.333333	6 --- 0.333333
7 --- 0.000000	7 --- 0.000000
10 --- 0.000000	10 --- 0.000000
Name: Survived, dtype: float64	Name: Survived, dtype: float64
IsAlone:	IsAlone:
0 --- 0.505650	0 --- 0.505650
1 --- 0.303538	1 --- 0.303538
Name: Survived, dtype: float64	Name: Survived, dtype: float64
Pclass:	Pclass:
1 --- 84.154687	1 --- 84.154687
2 --- 20.662183	2 --- 20.662183
3 --- 13.675550	3 --- 13.675550
Name: Fare, dtype: float64	Name: Fare, dtype: float64
Pclass:	Pclass:
1 --- 38.233441	1 --- 38.233441
2 --- 29.877638	2 --- 29.877638
3 --- 25.140620	3 --- 25.140620
Name: Age, dtype: float64	Name: Age, dtype: float64
Survived:	Survived:
0 --- 30.626179	0 --- 30.626179
1 --- 28.343690	1 --- 28.343690
Name: Age, dtype: float64	Name: Age, dtype: float64
Survived:	Survived:
0 --- 22.117887	0 --- 22.117887
1 --- 48.395488	1 --- 48.395488
Name: Fare, dtype: float64	Name: Fare, dtype: float64
1 --- 136	1 --- 136
3 --- 119	3 --- 119
2 --- 87	2 --- 87
Name: Pclass, dtype: int64	Name: Pclass, dtype: int64
3 --- 372	3 --- 372
2 --- 97	2 --- 97
1 --- 80	1 --- 80

4.2.8 Titanic Dataset Analysis and Data Cleaning-4

Problem Statement:

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

1. Get the number of survivors by gender (Sex).
2. Get the number of non-survivors by gender (Sex).
3. Get the number of survivors by embarkation location (Embarked_S).
4. Get the number of non-survivors by embarkation location (Embarked_S).
5. Calculate the percentage of children (Age < 18) who survived.
6. Calculate the percentage of adults (Age >= 18) who survived.
7. Get the median age of survivors.
8. Get the median age of non-survivors.
9. Get the median fare of survivors.
10. Get the median fare of non-survivors.

Code:

```
import pandas as pd
import numpy as np
```

```
# Load the Titanic dataset
```

```
data = pd.read_csv('Titanic-Dataset.csv')
data = pd.get_dummies(data, columns=['Embarked'], drop_first=True)
```

```
# 1. Get the number of survivors by gender
```

```
survivors_by_gender = data[data['Survived'] == 1]['Sex'].value_counts()
print(survivors_by_gender)
```

```
# 2. Get the number of non-survivors by gender
```

```
non_survivors_by_gender = data[data['Survived'] == 0]['Sex'].value_counts()
print(non_survivors_by_gender)
```

```
# 3. Get the number of survivors by embarked location
```

```
survivors_by_embarked_s = data[data['Survived'] == 1]['Embarked_S'].value_counts()
print(survivors_by_embarked_s)
```

```
# 4. Get the number of non-survivors by embarked location
```

```
non_survivors_by_embarked_s = data[data['Survived'] == 0]['Embarked_S'].value_counts()
```

```
print(non_survivors_by_embarked_s)
```

```
# 5. Calculate the percentage of children (Age < 18) who survived
```

```
children = data[data['Age'] < 18]
```

```
children_survival_rate = children['Survived'].mean()
```

```
print(children_survival_rate)
```

```
# 6. Calculate the percentage of adults (Age >= 18) who survived
```

```
adults = data[data['Age'] >= 18]
```

```
adults_survival_rate = adults['Survived'].mean()
```

```
print(adults_survival_rate)
```

```
# 7. Get the median age of survivors
```

```
median_age_survivors = data[data['Survived'] == 1]['Age'].median()
```

```
print(median_age_survivors)
```

```
# 8. Get the median age of non-survivors
```

```
median_age_non_survivors = data[data['Survived'] == 0]['Age'].median()
```

```
print(median_age_non_survivors)
```

```
# 9. Get the median fare of survivors
```

```
median_fare_survivors = data[data['Survived'] == 1]['Fare'].median()
```

```
print(median_fare_survivors)
```

```
# 10. Get the median fare of non-survivors
```

```
median_fare_non_survivors = data[data['Survived'] == 0]['Fare'].median()
```

```
print(median_fare_non_survivors)
```

Output:

Average time
0.030 s
30.00 ms

Maximum time
0.030 s
30.00 ms

1 out of 1 shown test case(s) passed

Test case 1 30 ms Debug

Expected output

Actual output

female....233	female....233
male.....109	male.....109
Name: Sex, dtype: int64	Name: Sex, dtype: int64
male.....468	male.....468
female....81	female....81
Name: Sex, dtype: int64	Name: Sex, dtype: int64
1...217	1...217
0...125	0...125
Name: Embarked_S, dtype: int64	Name: Embarked_S, dtype: int64
1...427	1...427
0...122	0...122
Name: Embarked_S, dtype: int64	Name: Embarked_S, dtype: int64
0.5398230088495575	0.5398230088495575
0.3810316139767055	0.3810316139767055
28.0	28.0
28.0	28.0
26.0	26.0
10.5	10.5